# Koushik C.

MBA (Operations), Business Analytics Minor — Targeting Operations & Analyst Roles

Bengaluru, India | +91 82772 73661 | chunchukoushik232@gmail.com | linkedin.com/in/koushik-c | github.com/koushik2302

## EDUCATION

**Master of Business Administration** Sep 2025 – Sep 2027

*PES University, Bengaluru — Major: Operations, Minor: Business Analytics*

CGPA: 8.09/10. Coursework: Production & Operations Management, Supply Chain, Business Analytics, Statistics, Managerial Economics, Financial Accounting.

**B.E., Computer Science & Engineering** 2020 – 2024

*Nitte Meenakshi Institute of Technology, Bengaluru*

## EXPERIENCE

**Storage Engineer** Oct 2024 – Jul 2025

*Hewlett Packard Enterprise (HPE)*

- Ran an L1/L2 incident queue for 50+ enterprise customers worldwide — triaged by severity and business impact, closed within SLA across time zones.
- Handled 20+ tickets at once during peak load, juggling prioritisation, escalation routing, and stakeholder updates in parallel.
- Chased actual root causes on recurring infrastructure failures instead of just patching symptoms, which brought down the repeat-ticket rate.

## PROJECTS

**Service Queue Simulation — Staffing, SLA Policy & the Utilisation Cliff** 2026

*Independent · Python (SimPy), JavaScript · github.com/koushik2302/queue-sim*

- Built a discrete-event simulation of a multi-tier support queue (non-stationary Poisson arrivals, lognormal service times, 4 severity tiers, 4 scheduling policies) to find the point where SLA compliance actually falls apart.
- Found that priority scheduling protects critical tickets while quietly wrecking the low-priority ones — at 95% utilisation, S1 waits stayed around 0.14h while S4's p95 wait hit 6.3 hours, even though the SLA dashboard still read 97% green. Recommended tracking p95 wait per tier instead of one aggregate pass/fail number.
- Validated the engine against the Erlang-C closed form before drawing any conclusions (23/24 checks inside the 95% CI), then stress-tested parameters at ±30% to separate real findings from artifacts of the inputs.
- Also built an interactive browser dashboard (JavaScript, Chart.js) so anyone can rerun the experiments and the validation suite themselves, rather than take the results on trust.

**Predicting Mental-Health Treatment Seeking in Tech** 2024

*Capstone · Python, scikit-learn, Flask · github.com/koushik2302/mental-health-prediction*

- Trained and compared SVM, Random Forest, Decision Tree, and a stacked ensemble on 1,259 survey responses (26 features) to predict treatment-seeking behaviour, then deployed it behind a small Flask app.
- The stacked ensemble (0.83) didn't beat a plain SVM baseline (0.85) — reported that honestly instead of burying it, since it's the more useful result.

## SKILLS

| | |
|---|---|
| **Operations** | Process improvement, root-cause analysis, capacity & staffing planning, SLA and incident management, queueing theory, supply chain fundamentals |
| **Analytics** | Discrete-event simulation, regression, forecasting, hypothesis testing, sensitivity analysis, A/B reasoning |
| **Technical** | Python (pandas, NumPy, SciPy, SimPy, scikit-learn), SQL, Advanced Excel, JavaScript, Git |
| **Tools** | Jira, Confluence, ServiceNow, Matplotlib, Chart.js |

## CERTIFICATIONS

- Lean Six Sigma: Green Belt Fundamentals — Alison (CPD Certified), 2026
- Enterprise Risk Management, Level 1 Course — Institute of Risk Management (course completion), 2026
- Advanced Software Engineering Job Simulation — Walmart USA via Forage